

# Assignment 2: Regression and Classification

AIGS/CSED515 Machine Learning

Instructor: Jungseul Ok

jungseul@postech.ac.kr

Due: 20:00pm Oct 22, 2020

## Remarks

- Group study and open discussion via LMS board are encouraged, however, assignment that your hand-in must be **of your own work**, and **hand-written**.
- Submit a scanned copy of your answer on LMS online in **a single PDF file**.
- Delayed submission may get some penalty in score: 5% off for delay of 0 ~ 4 hours; 20% off for delay of 4 ~ 24 hours; and delay longer than 24 hours will not be accepted.

1. [20pt; Linear Regression] We are given a dataset  $\mathcal{D} = \{(1, 1), (2, 1)\}$  containing two pairs  $(x, y)$  with  $x \in \mathbb{R}$  and  $y \in \mathbb{R}$ . We want to find the parameters  $\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \in \mathbb{R}^2$  of a linear regression model  $\mathbf{y} = w_1x + w_2$  using

$$\min_{\mathbf{w}} \frac{1}{2} \sum_{(x,y) \in \mathcal{D}} \left( y - \mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix} \right)^2. \quad (1)$$

- (a) [2 pt] Plot the given dataset and find the optimal  $\mathbf{w}^*$  by inspection.

**sol)** Line passing through points  $(x, y) = (1, 1)$  and  $(x, y) = (2, 1)$ .

$$\mathbf{w}^* = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

- (b) [4 pt] Write down  $\mathbf{y} \in \mathbb{R}^2$  and  $\mathbf{X} \in \mathbb{R}^{2 \times 2}$  which makes the following optimization equivalent to (1):

$$\min_{\mathbf{w}} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 \quad (2)$$

**sol)**

$$\mathbf{y} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}, \mathbf{X} = \begin{bmatrix} 1 & 1 \\ 2 & 1 \end{bmatrix},$$

- (c) [3 pt] Derive the general analytical solution for (2). Also plug in the values for the given dataset  $\mathcal{D}$  and compute the solution numerically.

**sol)** Setting the derivative to zero, we have  $\mathbf{X}^\top \mathbf{X} \mathbf{w} - \mathbf{X}^\top \mathbf{y} = 0$ . Hence,

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

- (d) [3 pt] There are several ways to compute this solution via PyTorch. Read the docs for the functions `torch.lstsq`, `torch.solve`, `torch.inverse`. Use all three approaches when completing the file `LinearRegression.py` and verify your answer.

**sol)** See `sol_LinearRegression.py`.

- (e) [6 pt] We are now given a dataset  $\mathcal{D}' = \{(0, 0), (1, 1), (2, 1)\}$  of pairs  $(x, y)$  with  $x \in \mathbb{R}$  and  $y \in \mathbb{R}$ . We want to fit a quadratic model  $\hat{y} = w_1x^2 + w_2x + w_3$  using (2). Specify the dimensions of the matrix  $\mathbf{X}$  and the vector  $\mathbf{y}$ . Also write down explicitly the matrix and vector using the values in  $\mathcal{D}'$ . Find the optimal solution  $\mathbf{w}^*$  and draw it together with the dataset into a plot.

sol)

$$\mathbf{X} \in \mathbb{R}^{3 \times 3}, \quad \mathbf{y} \in \mathbb{R}^3.$$

$$\mathbf{y} = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}, \quad \mathbf{X} = \begin{bmatrix} 0 & 0 & 1 \\ 1 & 1 & 1 \\ 4 & 2 & 1 \end{bmatrix}.$$

$$\mathbf{w}^* = \begin{bmatrix} -0.5 \\ 1.5 \\ 0 \end{bmatrix}.$$

Parabola passing through all three points in the dataset  $\mathcal{D}'$ .

- (f) [2 pt] Specify  $\mathbf{y}$  and  $\mathbf{X}$  in `LinearRegression2.py` to verify your answer for Problem 1e.

sol) See `sol_LinearRegression2.py`.

```
X = torch.Tensor([[0,0,1],[1,1,1],[4,2,1]])
```

```
y = torch.Tensor([[0],[1],[1]])
```

2. [20pt; Binary Logistic Regression] We are given a dataset  $\mathcal{D} = \{(-1, -1), (1, 1), (2, 1)\}$  containing three pairs  $(x, y)$ , where each  $x \in \mathbb{R}$  denotes a real-valued point and  $y \in \{-1, +1\}$  is the point's class label.

Assuming the samples in the dataset  $\mathcal{D}$  to be i.i.d. and using maximum likelihood, we want to train a logistic regression model parameterized by  $\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \in \mathbb{R}^2$ :

$$p(y \mid x) = \frac{1}{1 + \exp\left(-y\mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix}\right)} \quad (3)$$

- (a) [1pt] Instead of maximizing the likelihood we commonly minimize the negative log-likelihood ( $p(\mathcal{D} \mid \mathbf{w})$ ). Write the objective for the model given in (3) (don't plug in the instances of  $\mathcal{D}$ . In other words, write an optimization problem like (1)).

**sol)**

$$\min_{\mathbf{w}} \sum_{(x,y) \in \mathcal{D}} \log \left( 1 + \exp \left( -y\mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix} \right) \right) .$$

- (b) [3pt] Compute the derivative of the negative log-likelihood objective which you specified in Problem 2a (don't plug in the instances of  $\mathcal{D}$ ). Sketch a simple gradient-descent algorithm using pseudo-code (use  $\mathbf{w}$  for the parameters,  $\alpha$  for the learning rate,  $f$  for the objective function, and  $\mathbf{g} = \nabla_{\mathbf{w}} f$  for the gradient).

**sol)**

$$\sum_{(x,y) \in \mathcal{D}} \frac{\exp \left( -y\mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix} \right)}{\log \left( 1 + \exp \left( -y\mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix} \right) \right)} \left( -y \begin{bmatrix} x \\ 1 \end{bmatrix} \right)$$

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \mathbf{g}$$

- (c) [5pt] Implement the algorithm by completing `LogisticRegression.py`. State the code that you implemented. What is the optimal solution  $\mathbf{w}^*$  that your program found?

**sol)** See `sol_LogisticRegression.py`.

$$\mathbf{w}^* = \begin{bmatrix} 4.2385 \\ 0.0408 \end{bmatrix}$$

- (d) [3pt] If the third datapoint (2,1) was instead of (10; 1), would this influence the bias  $\mathbf{w}_2$  much? How about if we had used linear regression to fit  $\mathcal{D}$  as opposed to logistic regression? Provide a reason for your answer.

**sol)** No, it wouldn't since such an *easy* example contributes little to loss. However, it can significantly influence to the solution of linear regression which uses L2-loss instead of Log-loss.

- (e) [3pt] Instead of manually deriving and implementing the gradient we now want to take advantage of PyTorch auto-differentiation. Investigate `LogisticRegression2.py` and complete the update step using the instance named `optimizer`. What code did you add? If you compare the result of `LogisticRegression.py` with that of `LogisticRegression2.py` after an equal number of iterations, what do you realize?

**sol)** See `sol_LogisticRegression2.py`. The results are identical to each other as we optimize the same loss from the same initialization.

- (f) [5pt] Instead of manually implementing the cost function, we now want to take advantage of available functions in PyTorch, specifically `torch.nn.BCEWithLogitsLoss` which expects targets to be  $y \in \{0, 1\}$ . Consequently, you need to translate dataset  $\mathcal{D}$  with  $y \in \{-1, 1\}$  to dataset  $\mathcal{D}'$  with  $y \in \{0, 1\}$ . Write the probabilities  $p(y = 1 | x)$ ,  $p(y = 0 | x)$  and  $p(y | x)$  if we use `torch.nn.BCEWithLogitsLoss`. Complete `LogisticRegression3.py` and compare the obtained result after 100 iterations to the one obtained in previous functions.

**sol)**

$$p(y = 1 | x) = \frac{1}{1 + \exp\left(-\mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix}\right)}$$

$$p(y = 0 | x) = 1 - p(y = 1 | x) = \frac{1}{1 + \exp\left(\mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix}\right)},$$

which is equivalent to

$$p(y | x) = \frac{1}{1 + \exp\left((-2y + 1)\mathbf{w}^\top \begin{bmatrix} x \\ 1 \end{bmatrix}\right)}.$$

See `sol_LogisticRegression3.py`. The result is identical to the one obtained before.

3. [29 pt; Support Vector Machine]

We are given a dataset  $\mathcal{D} = \{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)}) : i = 1, 2, 3, 4\}$  of  $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}) = \left(\begin{bmatrix} 1 \\ 0 \end{bmatrix}, 1\right)$ ,  $(\mathbf{x}^{(2)}, \mathbf{y}^{(2)}) = \left(\begin{bmatrix} 0 \\ 1 \end{bmatrix}, 1\right)$ ,  $(\mathbf{x}^{(3)}, \mathbf{y}^{(3)}) = \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, -1\right)$ ,  $(\mathbf{x}^{(4)}, \mathbf{y}^{(4)}) = \left(\begin{bmatrix} -1 \\ -1 \end{bmatrix}, -1\right)$ . We want to train the parameters  $\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} \in \mathbb{R}^2$  and the bias  $b \in \mathbb{R}$  of a max-margin support vector machine (SVM) using: (for hyperparameter  $C > 0$ )

$$\min_{\mathbf{w}, b} \frac{C}{2} \|\mathbf{w}\|_2^2 \quad (4a)$$

$$\text{s.t. } \mathbf{y}^{(i)}(\mathbf{w}^\top \mathbf{x}^{(i)} + b) \geq 1 \quad \forall i = 1, 2, 3, 4. \quad (4b)$$

- (a) [5 pt] For the given data  $\mathcal{D}$ , how many constraints are part of the program in (4)? Specify all of them explicitly.

**sol)** Four constraints:

$$\begin{aligned} w_1 + b &\geq 1 \\ w_2 + b &\geq 1 \\ -b &\geq 1 \\ w_1 + w_2 - b &\geq 1. \end{aligned}$$

- (b) [8 pt] For  $b = 0$ ,  $b = -1$ , and  $b = -2$ , respectively, find the corresponding the set of feasible  $\mathbf{w}$  and optimal  $\mathbf{w}^*$  (if exists). Given only the three options  $b \in \{0, -1, -2\}$ , what is the optimal solution? Discuss whether a better solution exists.

**sol)**

When  $b = 0$ , the feasible set is empty, and thus no optimal solution exists.

When  $b = -1$ ,  $w_1 \geq 2$ ,  $w_2 \geq 2$  and  $\mathbf{w}^* = [2, 2]^\top$ .

When  $b = -2$ ,  $w_1 \geq 3$ ,  $w_2 \geq 3$  and  $\mathbf{w}^* = [3, 3]^\top$ .

Optimal solution is  $w_1 = 2 = w_2$  and  $b = -1$ . It is not possible to have a solution better than this since the feasible set is empty for  $b > -1$  and the cost function increases as  $b$  decreases from  $-1$ .

- (c) [5 pt] Draw the dataset in  $(x_1, x_2)$ -space using crosses ( $\times$ ) for the points belonging to class 1 and circles ( $\circ$ ) for the points belonging to class  $-1$ . Using your drawing, find the support vectors. Noting that those points for which the constraints hold with equality at the optimal solution, solve the resulting linear system w.r.t.  $\mathbf{w}$  and  $b$  and draw the solution into  $(x_1, x_2)$ -space.

**sol)** Support vectors are  $x^{(1)}$ ,  $x^{(2)}$  and  $x^{(3)}$ . The linear system consists of

$$w_1 + b = 1, w_2 + b = 1, -b = 1$$

of which solution is  $\mathbf{w} = [2, 2]^\top$  and  $b = -1$ .

- (d) [1 pt] What conditions do the datapoints have to fulfill such that the program in (4) has a feasible solution?

**sol)** Linearly separable.

- (e) [6 pt] In practice, for large datasets, it is hard to find the support vectors by inspection. A gradient based method is applicable. Using general notation, i.e., no plugging in  $\mathcal{D}$ , and introducing slack variables  $\boldsymbol{\zeta} = (\zeta_i)_{i=1,\dots,4}$  into (4), state the soft-margin problem with  $L_1$  penalty on  $\boldsymbol{\zeta}$  (including all constraints). Subsequently, reformulate this program into an unconstrained program. Finally obtain the gradient of this unconstrained program w.r.t.  $\mathbf{w}$  (use  $\frac{\partial}{\partial x} \max\{0, x\} = \mathbb{1}[x > 0]$ ). Compute the gradient at  $w_1 = 2$ ,  $w_2 = 2$  and  $b = -1$ , and discuss the impact of  $C$  and the relation between the max-margin and soft-margin SVMs.

**sol)**

The soft-margin problem with  $L_1$  penalty is obtained as follows:

$$\begin{aligned} \min_{\mathbf{w}, b, \boldsymbol{\zeta} \succeq 0} \quad & \frac{C}{2} \|\mathbf{w}\|_2^2 + \sum_i \zeta_i \\ \text{s.t.} \quad & y^{(i)}(\mathbf{w}^\top x^{(i)} + b) \geq 1 - \zeta_i. \end{aligned}$$

This can be reformulated as follows:

$$\min_{\mathbf{w}, b} f(\mathbf{w}, b) := \frac{C}{2} \|\mathbf{w}\|_2^2 + \sum_i \max\{0, 1 - y^{(i)}(\mathbf{w}^\top x^{(i)} + b)\}.$$

The gradient is obtained as:

$$\begin{aligned} \nabla_{\mathbf{w}} f &= C\mathbf{w} - \sum_i \mathbb{1}[1 - y^{(i)}(\mathbf{w}^\top x^{(i)} + b) > 0] (y^{(i)} x^{(i)}) \\ \frac{\partial f}{\partial b} &= - \sum_i \mathbb{1}[1 - y^{(i)}(\mathbf{w}^\top x^{(i)} + b) > 0] y^{(i)} \end{aligned}$$

When  $w_1 = 2$ ,  $w_2 = 2$  and  $b = -1$ ,  $\nabla_{\mathbf{w}} f = [2C, 2C]^\top$  and  $\frac{\partial f}{\partial b} = 0$ .

The choice of  $w_1 = 2$ ,  $w_2 = 2$  and  $b = -1$  is not optimal for  $C > 0$  as the gradient is non-zero for  $C > 0$ . However, as  $C$  goes to 0, the gradient converges to 0 at the choice of  $w_1 = 2$ ,  $w_2 = 2$  and  $b = -1$ . Hence, when we have sufficiently small  $C$ , the soft-margin approximates the max-margin.

- (f) [4 pt] Complete `SVM.py` with  $C = 1$  and verify your reply for the previous answer. What is the optimal solution  $(\mathbf{w}, b)$  that your program found and what is the corresponding loss? Explain the solution and what you observe when running the program, as well as how to fix this issue.

**sol)** See `sol_SVM.py`. Pytorch uses  $\frac{\partial}{\partial x} \max\{0, x\} = \mathbb{1}[x \geq 0]$  instead of  $\frac{\partial}{\partial x} \max\{0, x\} = \mathbb{1}[x > 0]$ . However, still we can check that  $\mathbf{w} = [2, 2]^\top$  and  $b = -1$  are not optimal. The program finds:

$$\mathbf{w} = [0.6674, 0.6674]^\top, \quad \text{and} \quad 0.3330$$

of which loss is about 1.779. The value of  $C = 1$  is too large, hence, we need to decrease it for small task loss.